

Cloud Algorithms and Mathematics

Barry Denby

Griffith College Dublin

April 30, 2020

Cloud Algorithms and Mathematics

- ▶ In this course thus far we have looked at the general theory of cloud computing systems
- ▶ We have covered areas as diverse as virtualisation, storage, security, and code optimisation
- ▶ In this lecture and the following lectures we will go deeper into the theory and explore some definitions and mathematics of cloud computing
- ▶ Implemented in most cloud systems under the hood
- ▶ Most of this material will be heavily based off Cloud Computing: Theory and Practice

Cloud Algorithms and Mathematics

- ▶ We will cover material related to
 - ▶ Parallelism in the cloud
 - ▶ Scheduling
 - ▶ Resource Management
- ▶ While this looks scattered there is a lot of material to choose from

Speedup of parallel computations

- ▶ Cloud computing environments are inherently parallel in nature
- ▶ Large problem sets can be run on the cloud with multiple virtual machines each processing a subset of the data
- ▶ Allows you to solve problems that demand more resources than those provided by a single system
- ▶ In the general case the speedup $S(N)$ of a parallel solution is $S(N) = \frac{T(1)}{T(N)}$

Speedup of parallel computations

- ▶ $T(1)$ is the execution time of the single threaded solution
- ▶ $T(N)$ is the execution time of the parallel solution
- ▶ Where N is the number of participating nodes in the parallel computations
- ▶ This is an important value that tells you both how scalable your algorithm is and how much performance it is generating

Speedup of parallel computations

- ▶ If $S(N) = 1$ then it means the parallel solution is no faster than the single threaded solution
- ▶ If $S(N) < 1$ then it means the parallel solution performs worse than the single threaded solution
- ▶ if $S(N) > 1$ then it means the parallel solution performs quicker than the single threaded solution

Speedup of parallel computations

- ▶ If $S(N) = N$ then we have linear speed up
 - ▶ There is no performance loss in the parallel solution for each node added
- ▶ If $S(N) < N$ then we have sub-linear speed up
 - ▶ There is some performance loss when a new node is added
 - ▶ The most common case as we expect some performance loss with additional synchronisation and communication
- ▶ If $S(N) > N$ then we have super-linear speed up
 - ▶ There is extra performance gained by running in parallel
 - ▶ Ideal but this is a very rare occurrence

Amhdal's law

- ▶ Gives the potential speed up of a parallel computation
- ▶ Defines a percentage of the computation that cannot be parallelised α
- ▶ Law is expressed as $S = \frac{1}{\alpha}$
- ▶ Where S is the speedup and α is the fraction of code that cannot be run in parallel
- ▶ As alpha increases then potential speedup will decrease and vice versa

Amdhal's law

- ▶ To prove this we define (σ) to be the sequential segment and (π) to be the parallel segment. As a result we can rewrite portions of our formulas as follows
- ▶ $T(1) = \sigma + \pi$, $T(N) = \sigma + \frac{\pi}{N}$, and $\alpha = \frac{\sigma}{\sigma+\pi}$
- ▶ By pluggins these into Amdhal's law and working with the equations we get (for large values of N)
- ▶ $S = \frac{1 + \frac{(1-\alpha)}{\alpha}}{1 + \frac{(1-\alpha)}{N\alpha}} = \frac{1}{\alpha + \frac{(1-\alpha)}{N}} \approx \frac{1}{\alpha}$

Gustafson's Law

- ▶ The problem with Amdahl's law is that it deals with a fixed problem size
- ▶ Realistically this will never be the case because, as more nodes are added to a parallel system there will be bigger datasets added to ensure speedup.
- ▶ Thus Gustafson modified Amdahl's law to account for a variable problem size.
- ▶ Gustafson's law is written as follows
- ▶ $S(N) = N - \alpha(N - 1)$

Gustafson's Law

- ▶ To prove this law we set the following
- ▶ $T(1) = \sigma + N\pi$, and $T(N) = \sigma + \pi$
- ▶ and plugging into the formula we get
- ▶ $S(N) = \frac{T(1)}{T(N)} = \frac{\sigma+N\pi}{\sigma+\pi} = \frac{\sigma}{\sigma+\pi} + \frac{\pi}{\sigma+\pi}$
- ▶ $= \alpha + N(1 - \alpha) = N - \alpha(N - 1)$

Coffman Conditions for concurrent deadlock

- ▶ In a system like the cloud there is the possibility with shared resources that deadlock can occur in the system
- ▶ Particularly if those shared resources are protected by locks to ensure serial access
- ▶ For deadlock to occur in a concurrent system there are four conditions that must be satisfied simultaneously
- ▶ These are known as the Coffman conditions

Coffman Conditions for concurrent deadlock

- ▶ Mutual Exclusion: At least one resource cannot be shared, only one process can access it at a time
- ▶ Hold and Wait: At least one process holds a resource and is waiting for other resources
- ▶ No preemption: The monitor process cannot force a process to release a resource
- ▶ Circular wait: P1 waits for P2 which waits for P3 waits for P4 ... PN waits for P1

Resource Management

- ▶ Resource management is core to any system as it affects the three basic criteria for evaluation of a system
 - ▶ Performance
 - ▶ Functionality
 - ▶ Cost
- ▶ Some functions may be too expensive or exhibit poor performance
- ▶ The strategies for resource management differ for all three cloud models

Resource Management

- ▶ Some problems like fluctuating workloads are shared by all three models
- ▶ Sometimes this can be predicted in advance e.g. seasonal loads
- ▶ Other times it will be unpredictable
- ▶ Requires spare available capacity
- ▶ And a monitoring system that can allocate resources to applications with little to no delay

Resource Management

- ▶ Unfortunately due to the size of the cloud, centralised control will not provide continuous service or performance guarantees
 - ▶ Great interest in autonomic policies
 - ▶ The idea being to get the system to manage itself
- ▶ There are many different policy areas that must be enforced and policed in a cloud system
- ▶ We will only scratch the surface and introduce some basic ways of determining who should get access to resources

Policies for Resource Management

- ▶ Cloud resource management policies can be loosely grouped into five classes
 - ▶ Admission Control
 - ▶ Capacity Allocation
 - ▶ Load Balancing
 - ▶ Energy Optimisation
 - ▶ Quality of Service
- ▶ Load Balancing and Energy Optimisation were talked about in much detail in an earlier lecture.

Admission Control

- ▶ The explicit goal of an admission control policy is to prevent the system from accepting workloads in violation of high-level system policies
- ▶ Not accepting work that would prevent current work from being completed
- ▶ Requires knowledge of the system global state
- ▶ Which in a system of cloud scale is obsolete at best

Capacity Allocation

- ▶ Allocation of resources to individual instances
- ▶ Locating resources is subject to multiple global optimisation constraints
- ▶ Requires a search of a very large search space
- ▶ However, individual system state can change very rapidly

Quality of Service

- ▶ Gives guarantees about the performance of all components in the cloud
- ▶ These guarantees should never be broken
- ▶ However, this is the most difficult to address (as we will see with scheduling later)
- ▶ But this is the most critical to cloud computing

Policies for Resource Management

- ▶ The problem here is that the already existing optimal solutions for these problems do not scale up to large systems
- ▶ They also only target single parts of resource management
- ▶ Also many have complex calculations that are too slow for a cloud to compute fully and in time
- ▶ Cloud needs algorithms that will address multiple areas at once

Mechanisms for Resource Control

- ▶ There are four basic implementations of resource management
- ▶ Control Theory: uses sensors and feedback to guarantee system stability and predict transient behaviour
 - ▶ However, this can only predict local not global behaviour
- ▶ Machine learning: a technique that could be applied to coordinate several autonomic system managers
 - ▶ Does not require a performance model of the system

Mechanisms for Resource Control

- ▶ Utility Based: requires a performance model and a mechanism to correlate user-level performance with cost
- ▶ Market-oriented/economic mechanisms:
 - ▶ Do not require a model of the system
 - ▶ Have auctions for bundles of resources
 - ▶ Simulates an economy for resources

Mechanisms for Resource Control

- ▶ Distinctions should be made between interactive and non interactive workloads
- ▶ The management techniques differ for both
- ▶ Interactive workloads need flow control and dynamic application placement
- ▶ Non-interactive workloads are focused on scheduling
- ▶ Trying to account for both types is a very difficult challenge

Scheduling Algorithms for Computing Clouds

- ▶ Scheduling is a critical component of cloud resource management
 - ▶ responsible for resource sharing and multiplexing on many levels
 - ▶ server → VMs → applications → threads etc
- ▶ Good scheduling algorithms must have the following properties
 - ▶ efficient, fair, and starvation-free
 - ▶ The objectives for the scheduler depend on the type of workloads submitted to the system the scheduler will manage


Scheduling Algorithms for Computing Clouds

- ▶ E.g. batch scheduler tries to maximise throughput and minimise turnaround time
 - ▶ Time between job submission and completion
- ▶ E.g. real-time system must meet deadlines and be predictable
 - ▶ There is a hard defined time limit for each task
- ▶ Schedules can either be preemptive or non-preemptive
 - ▶ Preemptive: may interrupt a low priority task to complete a high priority task
 - ▶ Non-preemptive: all tasks cannot be interrupted by the scheduler regardless of priority

Scheduling Algorithms for Computing Clouds

- ▶ Schedules must address two aspects of resource management
 - ▶ The quantity of resources that are allocated
 - ▶ When the resources are allocated
- ▶ Plotting these parameters on a graph leads to the figure on the next slide
 - ▶ Determines the different classes of schedules and their respective requirements
 - ▶ Best-effort, Soft-requirements, and Hard-requirements are the different classes

Scheduling Algorithms for Computing Clouds



Scheduling Algorithms for Computing Clouds

- ▶ Best-effort policies don't impose requirements on resource quantity or timing with respect to allocation of resources
- ▶ Soft-requirements require statistically guaranteed amounts of resources and timing constraints
 - ▶ They will be observed the majority of time but allows for cases where these guarantees can be broken
- ▶ Hard-requirements dictate that all limits on resource amounts and timing are respected
 - ▶ Nothing is allowed to break these limits

Fair Queuing

- ▶ Computing and communication are common operations on the cloud
 - ▶ Thus scheduling algorithms used in the cloud will try and link the two together
- ▶ Such scheduling algorithms must manage many quantities at the same time
 - ▶ Bandwidth: The amount of data each flow is allowed to transport
 - ▶ Timing: When the packets of individual flows are transmitted
 - ▶ Buffer space: Storage for packets allocated for each flow

Fair Queuing: FCFS

- ▶ The easiest solution to this is to use an FCFS scheduler
- ▶ However, it does not guarantee fairness as greedy flows can get more bandwidth
- ▶ Another option is to use a queue for each flow and service each flow in a round-robin manner
- ▶ However, this does not guarantee fairness of bandwidth allocation


Fair Queuing: Start-time Fair Queuing

- ▶ An algorithm for allocating CPU bandwidth originally designed for multimedia applications
- ▶ The idea is to organise CPU bandwidth consumers into a tree like structure
- ▶ Root node is the processor, leaves are individual threads
- ▶ A scheduler acts at each level of the hierarchy

Fair Queuing: Start-time Fair Queuing

- ▶ The fraction of processor bandwidth B_i allocated to the intermediate node i is the following
- ▶
$$\frac{B_i}{B} = \frac{w_i}{\sum_{j=1}^n w_j}$$
- ▶ Where w_i is the weight of tree node i and the summation of w_j is the combined weight of all nodes in the tree at the same level of that node
- ▶ An example of such a tree is shown on the next slide

Fair Queuing: Start-time Fair Queuing



Fair Queuing: Start-time Fair Queuing

- ▶ The Tree that you see here has two virtual machines with three applications
- ▶ Each VM and application weighted and scheduled by the node above
 - ▶ E.g. Server decides VM's weighting and VM's decide the weighting of individual applications
- ▶ The tree is used to calculate the proportion of CPU time that is allocated to each thread
 - ▶ Gives a cascading calculation for example $t(1,3)$ will get 1/3 of $A(1)$'s processing time which in turn gets 3/4 of $VM(1)$'s processor time
 - ▶ Roughly 5 to 6% of processor time

Fair Queuing: Start-time Fair Queuing

- ▶ When a VM is not active its bandwidth will be reallocated to other VMs that are active
- ▶ Similar for applications and threads
- ▶ If a thread is not runnable its allocation will be transferred to other threads